Unit Economics -- K Neurodsp

Anticloud FZ LLE * 0-1.gg * lois@0-1.gg * October 2026

Unit Economics -- K Neurodsp

Anticloud FZ LLE* 0-1.gg *lois@0-1.gg

Date: October 2026

Headline Metrics

Revenue Per Tier

Infrastructure Cost Model

Hardware: NVIDIA T4 16GB GPU

On-demand price: $0.40/hr (AWS g4dn, GCP T4, Lambda Cloud)

Throughput: PAX L5 Narrow L2 General 27B at ~5M tokens/hr on T4 (Q4 quantized, batch 8)

Cost per 1M tokens: $0.40 / 5 = $0.08

At scale (dedicated GPU, reserved):

* Reserved T4 instance: ~$0.18/hr
* Throughput: 5M tok/hr
* Cost per 1M tokens: **$0.036** (55% cheaper than on-demand)

Magic Number (Sales Efficiency)

Magic Number > 1.0 = healthy. >1.5 = exceptional. PAX projected at 1.64 due to low infrastructure cost, compliance-mandatory market (demand-pull), and no cloud dependency for delivery.

Burn Multiple

Burn multiple < 1.5x at scale = Sequoia-grade efficiency.

Cohort Retention Assumptions

Sensitivity Analysis (Cost per Token)

Margins improve at scale. Unit economics are hardware-driven, not human-driven.

Anticloud FZ LLE 0-1.gg *lois@0-1.gg* October 2026*

PAX L5 Narrow L2 General 27B $0.08/1M tokens on T4 GPU*

Harvard Dataverse DOI 10.7910/DVN/YMJKOG ORCID 0009-0009-2233-6107*